Website audit — playa-ai.org, July 2026

The definitive audit of the live site: claims adjudicated against the institutional record, plus structure, copy, tone, accessibility, SEO and machine readability. Includes the corrected account of the 80,000 figure — Burning Man's real attendance versus Playa AI's own participant counts — a full link and CTA inventory, a drift table crediting what the org fixed, and a prioritised fix list.

Prepared 2026-07-27 · **Site captured** 2026-07-27 · **Length** ~15,670 words
Evidence base 2,378 extracted records over 206 call transcripts, an 8.3 MB Mission Control export and 22 grant-work chats
Surfaces audited /, /work, /events, /team, /eelp, /living-trust-framework, /manifesto, connect.playai.party

Website audit — playa-ai.org, July 2026

Site captured and re-verified live on **2026-07-27**. Surfaces audited: `/`, `/work`, `/events`, `/team`, `/eelp`, `/living-trust-framework`, `/manifesto`, the undocumented routes `/privacy-toolkit` and `/classic/*`, and the external Connect app at `connect.playai.party`. Four lenses were run in parallel — claims, structure and IA, copy and tone, and accessibility/SEO/machine readability — and reconciled here. Claim inventory: `.vault-meta/site-claims.md`.

---

1. Executive summary

Verdict: not ready for a funder conversation today, but the gap is narrower than it looks and almost all of it is configuration and copy rather than rewriting. This site is more honest than its peer group. Eight workstreams each carry a status tier. `/work` publishes six unresolved problems under "Open questions". The status words on the Living Trust Framework and the EELP are *more conservative than what the organisation told the OpenAI Foundation eleven days earlier*. The manifesto states Burning Man's attendance correctly and attributes it to Burning Man. The copy discipline here is real and it should not be talked down.

What fails diligence is a short list, and every item on it fails on the site's *own* published standard rather than an imported one. The single deployed product records microphone audio without saying so, behind a privacy policy that returns a not-found page — while the org's own EELP principle 02 requires telling participants "what data it uses… before the experience begins." The flagship open-source artifact has no LICENSE file, so "open-source" is not yet legally true, and the grant application says it already is. Every one of the site's contact paths is a Gmail-compose URL, which means the donor path is broken right now for anyone on Outlook. And nothing on the site is server-rendered, so the careful, conservative, correct pages are invisible to exactly the crawlers, AI retrievers and automated diligence tools that read an organisation before a human does.

Fix those four and the site moves from "would not take the meeting" to "credible early-stage org with unusually honest labelling."

The numbers at a glance

Measure	Value
Routes audited	7 public + 6 undocumented (`/privacy-toolkit`, `/classic`, `/classic/events`, `/classic/team`, `/classic/work`, `/playa-ai-living-trust-framework`)
Total links across the site	111
Gmail-compose CTAs in rendered pages	29

Measure	Value
Gmail-compose URLs in the JS bundle	45
`mailto:` links anywhere	**0**
`<form>` elements anywhere	**0**
`<input>` / `<textarea>` elements anywhere	**0**
CTAs that are `<button>` with no href	7 (5 on the homepage alone)
Inbound links to `/manifesto` from anywhere on the site	**0**
Bytes served to a non-JS client, every route	1,717 (identical)
HTTP status for a nonsense path	200
Uncompressed payload	668 KB (132 KB with brotli)
Distinct portfolio lists the site publishes	4, overlapping by 3 of 10 items
Newest event on `/events`	2026-03-30 (119 days before capture)
Critical findings	6
Total findings	22

2. How to read this

The coverage rule

The brain ends before the site does. Mission Control runs 2026-02-22 to 2026-03-27; the grant chats are a single day, 2026-07-16; transcripts close 2026-03-25. The site is later than all of them. A difference between the site and the vault is therefore *usually* the org having moved on, not the org being wrong — those cases live in the drift table and are deliberately kept out of the findings so they are never misread as errors. What appears in section 3 is only what survives that test: claims the record shows are unsupported, internally inconsistent, or already known to the org to be wrong.

The numbers correction — read this before anything else

An earlier version of this audit framed the 80,000 figure as an unverified or disputed claim. **That framing was wrong and has been removed everywhere in this document.** The corrected account:

> **Burning Man genuinely has approximately 80,000 attendees.** That is a true, publicly verifiable fact about the *host event*. Nobody at Playa AI invented it, and it is entirely legitimate to cite as context for where this work is rooted.
>
> **Playa AI is a different thing and ran its own, separate, much smaller events.** Its own participant counts are in the hundreds.
>
> The error was never fabrication. **The error is category conflation** — presenting the host event's total attendance as if it were Playa AI's own participant, study, or deployment count.

The vault shows exactly this pattern. One source says "80,000+ **ATTENDEES**" — correct, that is Burning Man's population. Other instances silently convert the same number into "80,000+ **opt-in participants**" and "consented research participants" — incorrect, because those are Playa AI's own figures and they are far smaller. The vault page participant count: 80,000 vs 500 is best read not as a dispute about whether a number is true, but as the first detection of this conflation: 80,000 is the host event, ~500 is what one internal document gives for an actual Oracle deployment.

Two consequences follow, and both matter more than a retraction would:

1. **The recommendation is never "kill the 80,000 figure."** It is "say it accurately." Where outbound material reads "derived from 80,000 consented participants" or "field-tested at 80,000 participants," that sentence should be **rewritten**, not deleted — because the underlying fact (this work happens at an 80,000-person event) is real and worth stating.
2. **The website already gets this right, and that is the most important structural finding in this audit.** See finding 2.

---

3. Critical findings

Ordered by severity. Each is: what the site says verbatim, the evidence, why it matters, the specific fix, and an effort estimate.

Finding 1 — Connect records audio without disclosing it, and the privacy policy is a 404

Severity: critical.

Site (Connect pre-experience disclosure, `connect.playai.party`): the entire account of AI involvement is one clause — "*AI creates a memory you can keep.*" The words microphone, voice, audio and recording appear nowhere in the flow. The one link that could carry the burden, "*Read our privacy policy*", points to

`playa-ai.org/privacy`, which renders **"404 · PAGE NOT FOUND"** under the heading "*This page has left the map.*"

What the app actually does, from its own JavaScript bundle: `navigator.mediaDevices.getUserMedia({audio: true})` → `MediaRecorder` → chunks collected into an `audio/webm` Blob → `POST /api/connections`, which returns `transcript`, `intentionSummary`, `insights` and `oracleKeyword`. Voice is captured, transmitted, transcribed and run through inference.

Evidence from the record. The site's own `/work` page describes the correct behaviour — "*Microphone access requires explicit permission for each session*" and "*AI can synthesize the spoken conversation*" — so the org knows what the app does; the disclosure screen simply does not say it. The published EELP principle 02 at `/eelp` is unambiguous: "*Participants are told when AI is involved, what it does, what data it uses, and what the outputs may be — before the experience begins, in plain language.*" That principle exists to govern exactly this screen. The vault also carries an unresolved, explicitly blocking objection: does Playa AI have adequate technical/security staffing to build Connect records the argument that one person should not be sole coder of an app handling people's most vulnerable emotional data, and that a security audit was needed. No resolution appears anywhere. Note also that `/work` hedges where the homepage does not: "*deletion-first workflows explored*" versus the homepage's flat "*Private by default.*"

Compounding it: a stale archive page at the undocumented route `/privacy-toolkit` is live and self-canonicalising, titled "*Privacy Toolkit Archive — Playa AI*" and labelled "*PROTOCOL / OPEN TOOL · DRAFT / UNDER REVIEW*". It is currently the top candidate for any person or crawler searching "Playa AI privacy" — an archived draft standing in for a policy that does not exist.

Why it matters. This is the one artifact the org points to as proof it ships. An ethics-infrastructure organisation whose only deployed system violates its only published standard has a standard that reads as decoration. This is also the finding with legal edges: undisclosed audio capture, no privacy policy, no deletion route — while the same site repeats "*Deletion is a right, not a support ticket*" on three surfaces with nothing behind it.

Fix. (a) Rewrite the `/explain` screen to name the microphone, the recording, the upload, the transcription, the retention period and the deletion route, in plain language, before the button. (b) Ship `/privacy` — even one page saying what is collected, where it goes, how long it stays, and the address for a deletion request. Point the Connect app's link at it. (c) Wire an actual deletion path and name who executes it. (d) Add `<meta name="robots" content="noindex">` to `/privacy-toolkit` and `/classic/*`, or delete those routes. (e) Until (a)–(c) exist, do not cite the EELP as an operating condition of Connect.

Effort: policy drafting is the real cost — half a day of writing. The disclosure screen is one screen of copy. Routing and noindex, 1 hour.

Finding 2 — The site states the 80,000 figure correctly, and the architecture has hidden the page that does it

Severity: critical. This is the single most consequential structural fact on the site.

Site (`/manifesto`, badged "PUBLISHED · LIVING DOCUMENT"), verified verbatim in-browser 2026-07-27:

> "*Every year, 100 staff and 1,000 volunteers guide 80,000 people into a living experiment in radical self-reliance, gifting, decommodification, civic responsibility, and creative expression.*"

That is the correct construction, and it is worth naming precisely why. It states Burning Man's attendance. It attributes it to Burning Man. It spells out the operating ratio (100 staff, 1,000 volunteers) so no reader can mistake it for a study cohort. It never converts the number into Playa AI's own participants. **Nowhere on the live site is that conversion made.** The site's own participant number is stated separately and modestly, on both `/` and `/work`: "*EARLY PARTICIPANTS / 80+ / ACROSS PILOTS AND WORKSHOPS.*" The conflation that appears in outbound material — "*derived from 80,000 consented participants*", "*field-tested at 80,000 participants*" — **does not appear on the website at all.**

The website solved this problem in prose. Then the information architecture buried the solution.

Verified directly: `/manifesto` is absent from the nav. The header contains exactly four anchors (`Overview` → `/`, `Events` → `/events`, `Team` → `/team`, `Work` → `/work`). Enumerating every `<a>` element on `/`, `/work`, `/events`, `/team`, `/eelp` and `/living-trust-framework`, **not one links to** `/manifesto` — on the homepage, `document.querySelectorAll('a')` returns 32 links and zero with `href="/manifesto"`. All seven "Read the Manifesto" CTAs are `<button type="button">` elements with no href; clicking them opens a `[role=dialog]` modal with the URL unchanged, and the full text sits behind a second accordion click ("*In 1000 words / Read the full manifesto +*"). There is no sitemap: `robots.txt` is 23 bytes with no `Sitemap:` directive, and `/sitemap.xml` returns the SPA shell. The only way to reach `/manifesto` is to type the URL.

Why it matters, and this is the crux of the whole audit. The only citable, linkable, correct statement of the number is on a page nobody can reach and no draft author can copy from. That is not a footnote to the conflation — it is a *structural cause* of it. When the right sentence has no address, every new grant draft, cold email and policy submission reconstructs the number from memory, and memory drops the attribution. That is precisely how "80,000 attendees at the event we work at" becomes "80,000 consented participants in our study."

The highest-exposure instance is still open. Three outbound vectors carry the conflated phrasing with no recorded correction:

Vector	Status at corpus close	Record
Five funder outreach emails to named contacts, 2026-03-26/27, asserting EELP "*was derived from 80,000 consented participants in a high-complexity field environment (Burning Man)*"	believed sent; confirmation never recorded	March 2026 funder follow-ups
Mozilla cold email v1.2, subject line "*An open governance standard for experiential AI — field-tested at 80,000 participants*", send window 2026-04-01 to 04-03	**open**, never cancelled	Mozilla outreach subject-line decision, confirm v1.2 next-action

Vector	Status at corpus close	Record
OECD.AI Policy Navigator + Partnership on AI submission asserting "*a field study at Burning Man 2024 (80,000+ opt-in participants)*" and "*a two-year field study*"	draft copy only, pending approval	`raw/mission-control/items/documents/0013-oecdai-policy-navigator-eelp-submission-draft.json`

Two vectors are recorded as **not** in circulation and can be left alone: the PJMF concept note (`submitted: false` — PJMF concept note) and the LinkedIn About tagline, ready to implement with no record of going live (LinkedIn build package).

The OECD.AI / Partnership on AI submission is the highest-exposure item, because a policy-database entry that conflates the host event with the applicant's own cohort is very hard to unwind once indexed. **But the fix there is a rewrite for accuracy, not a retraction of a false claim** — the Burning Man context belongs in that submission; it just has to be labelled as context rather than as a participant count.

Fix. (a) Convert all seven manifesto `<button>` elements to `<a href="/manifesto">` and delete the modal. A modal cannot be cited, linked, shared or crawled, and this page's entire value is being citable. (b) Give `/manifesto` a nav home. (c) Lift the Burning Man paragraph into a short, permanent, linkable block — `/about#origin` or `/origin`, headed something like "*Where this work is rooted*" — so every future grant, submission and cold email has one URL to copy from instead of reconstructing the number. (d) Rewrite the three outbound vectors for accuracy: replace "*derived from 80,000 consented participants*" with the correct construction (see section 4), which keeps the Burning Man context and states Playa AI's own numbers as Playa AI's own. (e) File a decision record in the vault stating the distinction, so the conflation cannot return through a draft.

Effort: (a)–(b) 1 hour of front-end. (c) half a day. (d) 1–2 hours of rewriting per document. (e) 20 minutes. Giving the correct sentence an address is the cheapest available structural defence against this recurring.

Finding 3 — Every contact path on the site is a Gmail-compose link; there are no `mailto:` links and no forms

Severity: critical. This is currently breaking the donor path.

Site: "*REQUEST THE BRIEF*", "*CONTRIBUTE*", "*SCHEDULE DONOR CONVERSATION*", "*SUPPORT YEAR ONE →*" — all of them resolve to `https://mail.google.com/mail/?view=cm&fs=1&to=hello@playa-ai.org&su=…`

Evidence. Measured across the whole site: **111 links, 29 Gmail-compose CTAs, 0 `mailto:` links.** Grepping the production JS bundle finds **45** occurrences of `mail.google.com/mail/?view=cm` and `grep -oE 'mailto:[^"]{0,50}'` returns nothing at all. Verified live on the homepage: 32 links, 12 Gmail, 0 mailto, 0 `<form>` elements, 0 inputs.

The unique subject lines show exactly whose paths break: `Donor%20Inquiry`, `Donor%20Inquiry%20—%20Year%20One`, `Donor%20Brief%20Request`, `Founding%20Adopter%20Inquiry`, `Artist%20Grant%20Interest`, `Living%20Trust%20Framework%20Contribution`,

`Partner%20with%20Playa%20AI`, `Research%20Collaboration`, `Cultural%20Stewardship`.

Who is blocked: anyone on Outlook or Exchange — which is most institutional funders, foundations and universities; anyone in a corporate browser where Google is not signed in; anyone using a desktop mail client; anyone in a privacy browser or on a locked-down device; anyone in a jurisdiction where Google services are restricted. They do not get a friendly fallback. They hit a Google account wall on what looked like a contact button. It is also an accessibility defect: the link text gives no indication that activating it leaves the site for a third-party web application requiring separate authentication.

Why it matters. This is the highest-value five-minute fix on the site. Every donor, funder, Founding Adopter and artist path runs through it.

Fix. Replace all 45 with `mailto:hello@playa-ai.org?subject=…`, which respects whatever mail client the user actually has — Gmail included.

Effort: 30 minutes, find-and-replace.

Finding 4 — The Living Trust Framework repo has no LICENSE, so "open-source" is not yet true, and the grant says it already is

Severity: critical.

Site (`/`, `/work`, `/living-trust-framework`): "*WORKING DRAFT*" … "*intended for open-source release*" … "*The open-source license is under review.*" `/work` adds, unprompted, "*Technical practices listed here are implementation directions, not fully deployed capabilities.*"

The site is right. External check, 2026-07-27: `github.com/PlayaAI/playa-ai-living-trust-framework` is public, `"license": null`, created **2026-07-15** — the grant submission deadline — with 0 stars, 0 forks, 0 watchers, and 9 open issues all authored by the owner. Repo contents are only `README.md` (44,421 bytes) and `CONTRIBUTING.md` (1,562 bytes). The README is admirably candid: "*License status: Under review. No open-source or content license is currently granted unless and until an approved* `LICENSE` *file is added.*" `CONTRIBUTING.md` adds that contributors "*should not assume that submitted or existing material carries an open-source or content license.*" An unlicensed public repository is All Rights Reserved by default: visible, not open-source, not legally reusable. The site's status words match the artifact exactly.

Also worth flagging: the GitHub API reports the owner `"type": "User"`, not Organization. The canonical home of a standard intended to outlive its authors is a personal account — a continuity risk independent of the licence question.

The grant application does not match. The OpenAI Foundation commitment record carries status "*outstanding — asserted as done, verification never recorded*": the application describes the framework as an already-built, already-published open-source governance standard on GitHub. The artifact page records three tenses in a single day — "already publicly released on GitHub", "now being publishing", and a drafting chat refusing to assert it at all ("*Use this sentence only if the framework and its open-source materials are currently accessible*"). The verification action (verify publication status before referencing it) was **still open** at corpus end, and the project page states plainly: "*The corpus contains no record that this verification happened.*"

Why it matters. A reviewer reads "already publicly released on GitHub", opens the repo, finds it created the day before submission with no LICENSE and zero engagement, then opens the website and reads "the open-source license is under review." Three of the applicant's own surfaces disagree about whether the flagship governance artifact is published — and the one verifiable fact, no licence, means the word "open-source" is not yet true. It runs directly against the site's own Gifting principle, "*Open-source the standards. Share the learning publicly. The work becomes infrastructure when the lessons travel freely,*" and against §09's "*Designed to be forked, not just read*" on a repo that cannot legally be forked.

Fix. Add the LICENSE file — CC BY-SA 4.0 as the README's own "Historical License Proposal" suggests, or Apache-2.0 / CC0 if maximum adoption is the goal. That single commit makes the site's claim, the repo and the grant all true at once. If a licence genuinely cannot be chosen yet, send the OpenAI Foundation a one-line correction: public working draft, licence pending. Separately, convert `PlayaAI` to a GitHub Organization and transfer the repo.

Effort: 1 hour for the licence once the decision is made; 30 minutes for the org conversion. This is the single cheapest action on the entire list that turns a stated principle into a true one.

Finding 5 — "Privacy Toolkit in review" is the artifact name the org's own review struck

Severity: critical.

Site (`/` , §06 Tools), verified verbatim: "*Connect App released. Ikigai App piloted in Miami. Agent Labs running monthly. Privacy Toolkit in review.*"

What the record shows. Named real published tools over ChatGPT's invented Privacy Toolkit, 2026-07-16, states in terms: "*the 'Privacy Toolkit' referenced in an earlier ChatGPT draft does not correspond to a real published Playa AI artifact*", with the evidence line "*this definitively kills ChatGPT's fabricated 'Privacy Toolkit' — you don't need to invent one, you have two.*" The Living Trust Framework artifact page carries the same finding and warns that "*any downstream copy quoting the earlier draft is asserting a non-existent artifact.*"

The mitigating reading, stated fully. The Living Trust Framework project page records "privacy toolkit" as one of four genuine *internal aliases* for the Living Trust Framework, and "in review" is independently supported for that asset by the 2026-02-13 record of "privacy toolkit v1 in expert review" and by the decision approving "currently in expert review" as external copy. So the underlying consent work is real, and "in review" fairly describes it.

Why it still matters on either reading. "Privacy Toolkit" appears exactly once on the whole site — in the homepage tally — with no page, no link, no owner and no status badge. `/work` says "Eight workstreams" and enumerates them: Connect App, Ikigai App 1, Ikigai App 2, Relationship Nurturer, Living Trust Framework, EELP, Consent-First Data Trust, Burner-AI Manifesto. **No Privacy Toolkit.** So either it is the Living Trust Framework counted a second time under a second name with a second status word — which is padding, on a page whose next section says "*Credibility begins with honest status*" — or it is a tool named in a public tally that exists nowhere on the site and nowhere in the vault. A People-First AI Fund reviewer holding the grant file is looking at an artifact name the applicant's own review removed eleven days before this page was captured.

Fix. Delete the phrase from the homepage. If the consent toolkit is a real deliverable distinct from the Living Trust Framework, give it a `/work` entry, an owner and a status badge like the other eight.

Effort: one line of copy.

Finding 6 — Nothing is server-rendered, so the site's most careful pages are invisible to machines

Severity: critical. The mission-level argument for this is in section 9.

Evidence. Every route returns a byte-identical **1,717-byte** shell whose `<body>` is literally `<div id="root"></div>`. Diffed: `/` vs `/work` is identical; `/` vs `/nonsense-xyz` is identical. Bot user-agents were tested directly — Googlebot on `/manifesto`, GPTBot on `/work`, and `facebookexternalhit` on `/living-trust-framework` all receive the same 1,717 bytes. There is no prerendering tier for anyone.

Three defects compound:

1. **One social identity for the whole site.** The static shell hardcodes `<meta property="og:url" content="https://playa-ai.org/">` and `<meta property="og:title" content="Playa AI — AI Needs Culture.">` on *every* route. A funder sent a link to `/living-trust-framework` sees a preview card titled "Playa AI — AI Needs Culture." pointing at the homepage, in Slack, LinkedIn, X, WhatsApp and most email clients. The flagship framework has no shareable representation of its own. Platforms that canonicalise by `og:url` may collapse every share into one homepage entity.
2. **Soft-404s that self-canonicalise.** Every path returns HTTP 200 — verified on `/privacy`, `/this-path-does-not-exist-12345`, `/llms.txt`, `/ai.txt` and `/.well-known/security.txt`. Loading `/privacy` in a browser renders H1 "*This page has left the map.*" and title "*Page Not Found — Playa AI*" while simultaneously emitting `link rel=canonical` = `https://playa-ai.org/privacy`. A 404 page declaring itself canonical is an explicit instruction to index it. Meanwhile the *static* bytes any non-JS crawler receives for `/privacy` say `robots: index, follow` and title "Playa AI — AI Needs Culture." — which is worse than a plain 404. Link checkers and archive crawlers cannot detect any broken link on this site, because nothing ever signals broken.
3. **No sitemap and no structured data.** `/sitemap.xml` returns `content-type: text/html` and the SPA shell. `robots.txt` has no `Sitemap:` directive. There is no JSON-LD anywhere.

Why it matters. `/manifesto`, `/eelp` and `/living-trust-framework` — the three most careful, most conservative, most hedged pages the org has written — do not exist to any crawler, AI search index, archive tool or automated diligence system. The corrected framing of the numbers lives only in executed JavaScript, while the conflated framing sits in outbound PDFs and a policy-database submission as plain indexable text. **A retrieval system asking "how many participants does Playa AI have?" can read the error and cannot read the fix.** Fixing rendering is not a cosmetic SEO chore; it is how the accurate framing becomes reachable.

Fix. Prerender the seven real routes with their own `<title>`, `<meta name="description">`, `og:url` and `og:image` — for a Vite SPA this is `vite-plugin-prerender` or a static prerender step in CI; migrating to Next or Astro is the larger option but is not required. Prioritise `/`, `/work`, `/living-trust-framework`, `/manifesto`. Return a real HTTP 404 for unmatched paths (App Engine `error_handlers`), remove the

self-canonical from the 404 view and add `<meta name="robots" content="noindex">` to it. Generate `sitemap.xml` from the route table and add the `Sitemap:` line to `robots.txt`. Publish a plain-text `/llms.txt` that states Playa AI's own participant figures explicitly and separates them from Burning Man's attendance.

Effort: 1–2 days for prerendering; 3 hours for 404s and sitemap; 1 hour for `llms.txt`.

Finding 7 — `/get-involved` was scoped, specified and never built, which is why the site has zero forms

Severity: high.

Evidence. Verified in-browser: `document.querySelectorAll('form').length` = **0** and `main input, main textarea, main select` = **0**. Every conversion point on the entire site is a Gmail compose link.

The vault records that this was not the plan. Minimum viable launch scoped to 3 pages quotes the build package verbatim: "*Minimum viable launch: Homepage + /events + /get-involved*", with the stated rationale "*Homepage plus /events drives immediate action; /get-involved captures conversion.*" Two of the three shipped. `/get-involved` returns the 404 shell.

Earlier still, Single primary CTA (Get Involved), manifesto moved to page bottom records Alex Azzi agreeing with Massey Branscomb and Maria Gil that two competing CTAs reduce conversion, with the quoted evidence: "*we should be able to, you should be able to register without having to read the manifesto.*" The shipped site has **no registration mechanism of any kind**, and instead of one primary CTA the homepage carries twelve Gmail links.

Why it matters. The Gmail links are scar tissue from a missing page. Every other CTA defect on this site is downstream of `/get-involved` not existing.

Fix. Build it. One page, seven anchored sections (`#builders`, `#artists`, `#researchers`, `#funders`, `#stewards`, `#institutions`, `#general`), each with: who it is for, what you get, what is expected of you, what happens after you submit and by when — and a real form. That single page retires 29 Gmail links, gives the seven doors somewhere to land, and creates the first record of inbound interest the org can actually count.

Effort: 1–2 days including a form backend.

Finding 8 — You cannot sign the pledge

Severity: high.

Site: the "TAKE THE PLEDGE" CTA opens a modal headed "*FOUNDING SIGNATORIES / PUBLIC COMMITMENT*" / "*The Pledge*", inviting "*Sign it as an institution, a funder, or a person.*"

Evidence. Measured inside the modal: ten commitment statements, **0 forms, 0 inputs, 2 focusable elements** (a close button and one Gmail link), and **no signatory list**. The only mechanism to act on any of the ten commitments is a compose window with subject `Public%20Commitment%20—%20The%20Pledge`.

Why it matters. This is the highest-intent moment on the entire site — a funder or institution ready to publicly commit — and it terminates in a mail client. A signatory list would also be the indexable social proof the site conspicuously lacks.

Fix. Add a real signing form posting to a simple store, and render the resulting signatory list on the page. Interim: at minimum make the link a `mailto:` so it works everywhere.

Effort: 1 day for a real form; 15 minutes for the interim.

Finding 9 — A microgrant program is soliciting money with nothing behind it in the record

Severity: high.

Site (`/`), verified verbatim: "*…microgrant program is in development, and the artist interest list is open.*" The donor CTA on the same page: "*Fund the standards work, governance design, Connect pilots, artist microgrants, documentation, and public field reporting.*"

What the record shows. Nothing. A case-insensitive search for "microgrant", "micro-grant" and "micro grant" across the entire wiki returns zero hits. It is absent from the OpenAI Foundation 12-month plan, the seven-item Field Lab 2026 plan promised to that funder at the 2026-07-15 deadline — none of the seven items is regranting or artist grants. The nearest analogue is the Burning Man 2025 activation: "*The grant program received one submission… No grant disbursement appears anywhere in the corpus.*"

Compounding it, the artist CTA's Gmail subject line commits harder than its own label: the button reads "*Join the artist interest list*" and the subject sent is `Artist Grant Interest` — asserting a grant program the same page calls "in development."

Why it matters, even though absence from the vault is not proof of absence. This is a status word with a live financial ask attached, missing from the twelve-month plan given to the same funder the same month, from an org whose only recorded grant program disbursed nothing. Adjacent and unresolved: is MAF's "No" regranting answer accurate given its 2024 Form 990 grant payments records that the fiscal sponsor answered "No" to the application's regranting question despite reporting $28,609 in grants to individuals, calls it "*the largest diligence vulnerability*", and the written confirmation required to settle it is not recorded as obtained. A publicly advertised artist-grant program sitting next to that answer is the first question a diligence officer asks.

Fix. Either publish the program's eligibility, budget, review process and timeline, or take "artist microgrants" out of the donor CTA until it exists. Change the Gmail subject from `Artist Grant Interest` to `Artist Interest List`. Close the MAF regranting question in writing before either.

Effort: 15 minutes for the copy and subject line; the regranting confirmation is an email and a wait.

Finding 10 — The homepage understates the org roughly five-fold, using the site's own evidence

Severity: high.

Site (`/` , `/work`), verified verbatim: "*EARLY PARTICIPANTS / 80+ / ACROSS PILOTS AND WORKSHOPS.*"

What the record shows. 80+ is Connect alone. Portfolio of small shipped pilots lists them separately: "*Connect — a shipped pilot with 80+ users*" and "*Ikigai — an AI-facilitated purpose-discovery experience used in in-person workshops reaching roughly 80 people*", plus "*OpenClaw hatchathons… 358 community members from curiosity to capability*" and "*roughly 200 builders, artists, researchers, and funders meeting fortnightly.*" The work-page curation decision confirms Connect and Ikigai were selected as two items "*both with real participant counts.*"

Why it matters. This is a site-internal contradiction: `/work` attributes 80+ to Connect and ~80 to Ikigai on the same page where `/` presents 80+ as the org-wide total. An aggregate cannot equal its largest component. It reads as "your numbers don't add up" — and unusually, the error runs *against* the org. It is also the mirror image of the conflation problem: having correctly refused to borrow the host event's number, the site now under-counts its own.

Fix. State the real aggregate with a definition. See section 4 for the exact recommended wording. Before publishing 358, define it — cumulative attendees across N sessions since when — because that figure enters the record on a single working day of grant drafting with no attendance list behind it, which the canonical page itself flags.

Effort: 30 minutes of copy, plus whatever it takes to define 358.

Finding 11 — The event roster counts places the founder travelled to, and omits the event that created the format

Severity: medium-high.

Site (`/`), verified verbatim: "*9 public events across Dubai, Detroit, Lisbon, Davos, New York, St. Moritz, and online — 2025 to 2026.*" `/events` : "*UNIQUE EVENTS 9 · LOCATIONS 7 · YEARS 2025–26 · CORE FORMATS 4.*"

First, a correction to our own earlier inventory: **the arithmetic is fine.** Six cities plus online is seven locations. There is no six-versus-seven discrepancy and we should not report one.

What the record shows. `wiki/events/` records roughly six Playa AI convenings through corpus close — the London salon, the Lisbon Web Summit launch, European Ignition, the Feb 2026 Agentic Salon, Hackathon 1 New York (46 registered / 19 attended), and Hackathon 2. Nine by July is entirely plausible — four months of unrecorded activity sits between corpus close and capture, so the *count* is drift, not error. The **roster** is the finding:

- **Detroit** appears in the corpus only as Summit Series travel ("Summit, Davos, Burning Man, Cabo, Detroit" as personal non-negotiables). No Playa AI-hosted Detroit event exists in the vault. Summit Detroit runs in June, after transcript close, so a real June event is possible — unsupported, not contradicted.

- **St. Moritz** occurs once in the entire corpus, as a community member recalling meeting the founder socially. St. Moritz is a January-to-March destination, squarely inside the coverage window; a Playa AI public event there should be in the corpus and is not.
- **Dubai** is recorded as a networking dinner at which the founder pitched, not a hosted event.
- **Davos** goes the other way: the 2026 planning scaled the ambition *down* once venue rates of "$35,000 to $60,000 to hold a space for two hours" settled it, in favour of small self-run gatherings.
- **London** — the one event the vault treats as the origin of the org's signature format; the salon series says the format invented there "*became the template Playa AI reused in Lisbon and pitched at Davos*" — is **absent from the site entirely**.

Two of the nine events on `/events` are attributed to a host that is not Playa AI. That attribution question is in the private annex.

Why it matters. A funder who asks "which nine?" gets a list the record cannot reconcile, and the pattern it exposes — founder presence rendered as institutional activity — is a smaller instance of the same category error as the participant conflation: *proximity to a thing is not the same as having done the thing.* It is also self-defeating: the strongest event in the archive is the one that was left out.

Fix. Publish the event log. Nine rows: date, city, host, format, attendance, link. Define what counts as a Playa AI public event versus a venue the founder attended or a third party hosted. Put London back in.

Effort: half a day, mostly reconciliation.

Finding 12 — The present tense is four months stale

Severity: medium.

Site: `/events` header — "*PUBLIC EVENT ARCHIVE · SALONS / HACKATHONS / FIELD GATHERINGS · UPDATED CONTINUOUSLY*" — and the closing line "*The archive grows. The next room is already forming.*" The nine events, dates extracted in order: MAR 30 2026, MAR 18 2026, MAR 4 2026, FEB 6 2026, JAN 26 2026, JAN 23 2026, NOV 13 2025, JUN 4 2025, MAR 29 2025. **Newest is 2026-03-30; today is 2026-07-27.** Testing the page text for `upcoming|register|rsvp` returns false — there is no upcoming section at all. The homepage simultaneously calls Agent Labs **"ACTIVE"** and "*running monthly*", and the footer date-stamps the whole site "*SPRING 2026 · BLACK ROCK DESERT ORIGIN*".

What the record shows. The vault corroborates the hatchathon programme up to March 2026 (AI hackathon series) and records the 358-participant figure, but records **no monthly cadence**. The Data Trust is another case of momentum language over a stalled state: the site says "*Currently in governance design phase*", while the project page records status *stalled* — no Trust Council seat filled, the Council never convened, the Public Use Registry not live, the deletion pipeline not operational.

Why it matters. "Active", "running monthly" and "updated continuously" are checkable status words, not framing, so drift is not a defence. A claim that decays into falsehood by the passage of time is worse than no claim, because the reader can date it themselves.

Fix. Add an "Upcoming" block at the top of `/events` — even if it reads "Next session: TBA — join the list to hear first" with a real form. Link the Luma *calendar* so people can subscribe once instead of per event. Replace "UPDATED CONTINUOUSLY" with "Last updated: `<date>`", and put a `Last updated` stamp in the footer of every page. A dated quiet quarter reads as honesty; an undated "continuously" reads as abandonment.

Effort: 2 hours.

Finding 13 — `/manifesto` lists five unbuilt programs with no status labels, on a site whose discipline is status labels

Severity: medium.

Site (`/manifesto`, section "What We're Building"), verbatim: THE ORACLE — "*A living AI installation trained on the Playa's stories, art, and wisdom.*" THE FIRST AI CAMP. AI FOR HUMANITY — "*An open-source, community-governed foundation model trained on the spirit of the Playa.*" THE AI CODE OF THE PLAYA. A DIGITAL BURNING MAN — "*A permanent AI/VR-powered Playa experience.*" Present progressive, no status badge anywhere near them.

What the record shows. The manifesto is genuine and its status word is right: v1 was synthesised and read aloud on 2025-03-04 and used in outreach to Burning Man leadership and in the Eric Schmidt approach (2025 public launch and manifesto). The problem is that the aspiration list is not labelled the way `/work` labels everything else. An org whose measured output is one web app with 80 users is, on this page, telling funders it is building a foundation model and a persistent VR world. "Oracle" also reappears undefined in Living Trust §10 Phase 03 as a project the framework will be applied to — named as real, never explained.

Minor and separate: the page footer stamps the manifesto "*MANIFESTO · PLAYA AI · PUBLISHED 2026*", while the record dates drafting and first circulation to February–March 2025 and describes it in March 2026 as over a year old. Wrong year on a founding document.

Fix. Add one line above the section — "These are the directions we intend to build, not current programs" — or badge each item the way `/work` does. Fix the year stamp. And give the page a CTA: it is the most persuasive prose on the site and it currently converts nobody.

Effort: 30 minutes.

Finding 14 — Two phrases the org formally prohibited for public website copy are live, one of them in the hero

Severity: medium. This is a governance-of-content problem as much as a copy problem.

What the record says. Prohibited phrases for public-facing website copy, status `holds`: the rebuilt website must avoid "*Human-AI co-evolution,*" "*Born on the playa, gifted to the planet,*" "*Conscious Intelligence,*" and any pre-consent claims about how the Oracle "feels." Stated rationale: language like this

"reads as a spiritual TED talk and causes skeptical institutional audiences (Davos delegates, Stanford researchers) to bounce." The decision's own evidence line: *"If it sounds like a TED talk → rewrite. If it sounds like a governance document with a soul → ship it."*

What is live, verified 2026-07-27. The homepage hero reads **"Born on the Playa. Gifted to the Planet."** The manifesto block repeats **"Born on the Playa. Gifted to the planet."** The `/events` opening paragraph reads *"Playa AI events are live laboratories for builders, founders, researchers, artists, and operators exploring human-AI co-evolution with intention."* Two of the three prohibited phrases shipped, three occurrences, one of them the first line under the hero headline. ("Conscious Intelligence" is correctly absent.)

Why it matters. A locked decision with `status: holds` did not survive the build, and nothing in the architecture would have caught it. A decision that cannot be checked against the shipped artifact is not a decision.

Fix. Remove both phrases. Then close the loop structurally: keep the prohibited-phrase list in the repo next to the copy and grep it in CI.

Effort: 20 minutes for the copy; 1 hour for the CI check.

Finding 15 — The site enumerates its own portfolio four times and never the same way twice

Severity: medium-high.

Evidence. Homepage §03 *"CURRENT INITIATIVES · What is being built now"* lists **five**: Connect App (LIVE PILOT), Living Trust Framework (WORKING DRAFT), EELP (DRAFT), Human Connection Research / JOYSCORE EXPERIMENT (ACTIVE), Agent Labs / OPENCLAW HACKATHONS (ACTIVE). `/work` §02 lists **eight**: Connect App, Ikigai App 1, Ikigai App 2, Relationship Nurturer, Playa AI Living Trust Framework, EELP, Consent-First Data Trust, Burner-AI Manifesto. **The overlap is three items out of a combined ten.** JoyScore and Agent Labs are on the homepage and are not among the eight workstreams; Ikigai 1, Ikigai 2, Relationship Nurturer and the Data Trust are on `/work` and absent from the homepage. Homepage §06 then produces a **fourth** list — *"Connect App released. Ikigai App piloted in Miami. Agent Labs running monthly. Privacy Toolkit in review."* — naming a Privacy Toolkit that appears in neither of the other two. The homepage link *"SEE ALL WORKSTREAMS ON THE WORK PAGE →"* frames §03 as a subset of `/work`, which it is not.

Why it matters. The homepage is not a summary of `/work`; it is a competing portfolio. A reviewer cross-checking two pages of the same site finds two different organisations.

Fix. Pick one canonical list — the eight on `/work` — and make the homepage a strict, labelled subset: *"Three of eight workstreams. See all eight →"*. If JoyScore and Agent Labs are real workstreams they belong in the eight (making ten, and the `/work` counter changes with them). Delete the §06 prose tally; a stats strip and a portfolio list should never be two different portfolios.

Effort: 1 hour.

Finding 16 — Accessibility: motion and focus

Severity: high (both items).

`prefers-reduced-motion` **is ignored at both layers simultaneously.** Two independent checks, both negative. CSS: zero occurrences of `prefers-reduced-motion` in `index-BOG9hGSA.css`. JS: zero occurrences of `useReducedMotion` in the app bundle. The single `prefers-reduced-motion` string in the bundle is framer-motion's internal plumbing, and the decisive line is the library's context default — `createContext({transformPagePoint:a=>a,isStatic:!1,reducedMotion:"never"})` — with the branch `this.reducedMotionConfig==="never"?this.shouldReduceMotion=!1`. framer-motion defaults to "never", so the OS-level Reduce Motion setting is ignored unless the app opts in, and there is no `MotionConfig` override anywhere. Against that: **46** `whileInView` **usages and 55** `initial:"hidden"` **blocks** with variants `hidden:{opacity:0,y:20}`. A user with a vestibular disorder who has set Reduce Motion at the OS level still gets 46 slide-and-fade animations. Measured live on the homepage: **6,794 of 12,388 text characters (54.8%) sit at effective opacity below 0.1** until scrolled — including the H1 and the primary CTAs.

Fix: wrap the app root in `<MotionConfig reducedMotion="user">` — one line, makes framer-motion honour the OS setting globally. Add a CSS backstop: `@media (prefers-reduced-motion: reduce) { *,*::before,*::after { animation-duration:.01ms !important; transition-duration:.01ms !important; } }`. **Effort: 30 minutes.**

The pledge modal never moves focus into itself. `role="dialog"`, `aria-modal="true"` and `aria-labelledby="pledge-modal-title"` are all correctly set and the label resolves — but `document.activeElement` after opening is still `BODY`, and `dialog.contains(document.activeElement)` is `false`. There are **37 background focusables still tabbable**, `document.querySelectorAll('[inert]').length` is 0, and nothing outside is aria-hidden. Because `aria-modal="true"` is set, screen readers confine the virtual cursor to a dialog that focus never entered — a known stranding pattern. Body scroll locking is correctly applied.

Fix: on open, move focus to the dialog container or its first heading; on close, restore focus to the triggering button. Add a focus trap, or `inert` on the background wrapper. The ARIA scaffolding is already right, so this is genuinely the last mile. **Effort: 1 hour.**

Also: all images carry `alt=""`, marking content images as decorative. Audit and write real alt text for any image that carries meaning. **Effort: 1 hour.**

Finding 17 — Zero compression: 668 KB ships where 132 KB would do

Severity: high, and it is the best effort-to-benefit ratio on the site.

Evidence. Requested with `Accept-Encoding: br, gzip, deflate`; the response carries **no** `content-encoding` **header at all**. `content-length: 637938` for `/assets/index-DTO4XUBn.js`, `content-length: 44793` for the CSS. Compressed locally: gzip -9 → 176,221 bytes (72% saving); brotli -q 11 → 127,678 bytes (80% saving). CSS goes 44,793 → 6,838 bytes under brotli. Total page weight today is 668 KB; correctly compressed it is 132 KB. Compounding it, `cache-control: private` is set on content-hash-

fingerprinted immutable assets, defeating CDN caching and forcing revalidation — hashed filenames exist precisely so they can be cached for a year. The server is `Google Frontend` (App Engine static hosting), where this is a config setting, not a code change.

Fix. Enable brotli/gzip on the static handler; set `cache-control: public, max-age=31536000, immutable` on `/assets/*` and keep `no-cache` on `index.html` only. **Effort: 30 minutes, config file only.**

Finding 18 — Navigation: four links for seven pages, a dead decorative "nav item", and a Contact link that is `display:none` on mobile

Severity: high.

Full header markup, captured verbatim:

```
<nav class="site-header-nav flex items-center gap-5 md:gap-8">
  <a href="/">Overview</a> <a href="/events">Events</a>
  <a href="/team">Team</a> <a href="/work">Work</a>
</nav>
<div class="hidden md:flex items-center gap-4 shrink-0">
  <span class="text-muted-foreground hidden lg:inline">MMXXVI</span>
  <a href="/#contact">Contact</a>
</div>
```

Three consequences.

1. `MMXXVI` **is a** `<span>` **, not a link.** It carries the identical `text-muted-foreground` class as the unvisited nav anchors, so it reads as a fifth nav item that does nothing, and it vanishes entirely below 1024px. It is also opaque — nothing on the site glosses it as 2026.
2. `Contact` **resolves to** `/#contact`, which is the closing block headed "*THIS IS THE INVITATION / Will you help shape AI with human depth?*" It contains three CTAs (two Gmail, one modal button) and **no address, no phone, no form, no named person and no response commitment.** The word "contact" leads to a block with no contact information.
3. **The Contact button lives in a** `hidden md:flex` **container** and the header contains no `<button>` — there is no hamburger. Below 768px the only contact CTA in the chrome is `display:none`, with nothing to replace it. The vault records this exact class of bug being raised sixteen months ago: fix mobile UX on Playa AI website (hidden buttons, broken hover), raised by Maria Gil and Massey Branscomb, status still `open`.

Fix. Delete the `MMXXVI` span or make it a link. Point `Contact` at a real `/contact` (or `/get-involved#contact`) carrying an email address in text, the fiscal sponsor's legal name and link, a postal address if one exists, and a stated response window. Add a mobile menu, or move Contact inside `<nav>` so it survives below 768px. **Effort: 2 hours.**

Finding 19 — Every subpage is a leaf, and the back-link uses vocabulary the reader has never seen

Severity: medium.

Internal outbound links per page, excluding the global header: `/manifesto` — one ("← BACK TO DISPATCH" → `/`). `/team` — one (same). `/events` — one (same), plus nine outbound `luma.com` links. `/work` — three (`/living-trust-framework`, `/eelp`, `#work-index`) plus the back-link. `/eelp` — one (`/living-trust-framework`), the only lateral sibling link on the site. No breadcrumbs, no "related", no next/previous.

"Dispatch" appears nowhere else on the site — the homepage is called "Overview" in the nav and "dispatch" in every footer, so the back-link names a page the reader has never been told exists under that name.

Two dead ends stand out. `/manifesto` **carries no CTA at all** — no join, no pledge, no contact, nothing but the back-link. And `/work` **has no link to Connect**: no anchor on `/work` points at `connect.playai.party`; its Connect CTA is "Join pilot →", a Gmail link. You can read 1,100 words about the flagship pilot and not be able to open it.

Fix. Give every page a footer "Next" strip of two or three real onward links chosen for that reader. Rename "← BACK TO DISPATCH" to "← Overview" to match the nav. Add "Open Connect →" to the Connect card on `/work`. Put a single CTA at the foot of `/manifesto`. **Effort: 2 hours.**

Finding 20 — `/team` is the thinnest page on the site, lists all ten people twice, and there is no About page

Severity: medium.

`/team` is **2,355 characters** for an organisation whose credibility rests on its people and council — against 13,305 for `/work` and ~12,700 for the homepage. Every one of the ten names appears **twice**: once in the hero roster ("Alex Azzi FOUNDER & EXECUTIVE DIRECTOR / Tina Woods HEAD OF RESEARCH / …") and again in the numbered blocks ("001 · FOUNDER Alex Azzi", "A·002 · COUNCIL Richard Rothenberg", …). So real information density is roughly half of an already-thin page.

No bios, no affiliations, no prior roles, no outbound links. The entire biographical payload for the founder is "*HOLDS THE VISION. TENDS THE FLAME.*" and for the CTO "*PROGRAMS & OPERATIONS.*" The two renderings disagree with each other: Rob Griffitts is "LEGAL LEAD" in the roster and "LEGAL & POLICY" in the council block; Alvin appears as both "Alvin Graylin" and "Alvin Wang Graylin". Council numbering runs A·002 to A·007 with no A·001, which reads as an unrenumbered removal.

`/about` returns the 404 shell, and the vault explains why — pending: Founding Story approval for external use records the Founding Story drafted into `/about` and never approved, blocking both the page and the pitch deck. Meanwhile "MYSTIC ARTS FOUNDATION" is named as plain text in the footer of all seven pages and is **never a link** — the single most diligence-relevant external entity on the site cannot be clicked.

Fix. Merge `/team` into a real `/about`: origin, mission, fiscal sponsorship (with Mystic Arts Foundation linked and its EIN stated), legal structure, the people, and the Council Charter that is already written and genuinely good. Delete the duplicate roster. Add one line of affiliation and one outbound link per person — two council members have substantial verifiable institutional histories and this is the cheapest credibility on the site. Reconcile the two title variants and the two spellings of Alvin's name; restore or renumber A·001. **Effort:** half a day, plus unblocking the Founding Story approval.

Finding 21 — `/work`'s strongest asset is ten screens down, and the section labelled "Work Index" is not an index

Severity: medium.

`/work` §01 is titled "Work Index", subhead "*Eight workstreams. One public-interest system.*" What follows is a four-figure stats strip (WORKSTREAMS 8 / MODES 3 / EARLY PARTICIPANTS 80+ / LIVING MANIFESTO 1) and two paragraphs of prose. There is no index: no jump links, no list of the eight names, no filter, no status table. The reader must scroll linearly through eight long-form cards.

Meanwhile the section headed "*OPEN QUESTIONS / What remains unresolved*" — "*The honest edges matter. These are the questions the work must answer before it deserves to scale*", followed by six numbered items including "05 Security testing" and "04 Anti-sycophancy" — **begins at character 11,963 of a 13,305-character page.** It is the most credibility-generating content anywhere on the site and it is the last thing before the footer.

`/work` is also the page whose `<article>` markup causes generic content extractors to return only workstream 02 (Ikigai App 1) as the page's main content.

Fix. Make §01 an actual index: eight rows, each with name, mode, status badge, owner and an anchor jump. Move "What remains unresolved" up, directly beneath the index — an org that leads with its open problems is making an argument, and burying it forfeits the argument. Reserve `<article>` for the page-level container so extractors and AI readers get the whole page. **Effort: 2 hours.**

Finding 22 — The product surface is scattered across five domains, two of them personal prototyping hosts

Severity: low-medium.

Live outbound product links: Connect → `https://connect.playai.party` (linked from the homepage only; the vault records a decision to move off `playaiparty` to `playa-ai.org`). Ikigai App 1 → `https://ikigai-rmarji.replit.app`, labelled "*OPEN APP — CURRENTLY UNAVAILABLE →*" — a Replit URL containing an individual's username. Ikigai App 2 → `https://ikigai-ai-qgergqr8.manus.space`, labelled "Open prototype →" — a Manus scratch deployment with a random hash. Living Trust Framework → `github.com/PlayaAI/playa-ai-living-trust-framework`.

A funder or adopter following the work touches five hostnames, none of which is `playa-ai.org`, and two of which announce themselves as throwaway prototyping environments under a personal account.

To the site's credit, the one duplicate route it does have is handled correctly: `/playa-ai-living-trust-framework` serves the same page as `/living-trust-framework` and correctly sets `<link rel=canonical>` to the latter.

Fix. Put everything under `playa-ai.org` subdomains or subpaths — `connect.playa-ai.org`, `ikigai.playa-ai.org` — or at minimum a redirect. **Effort:** a few hours of DNS and config.

---

4. The numbers, explained properly

This is the section to read before writing any external document.

The two categories

Figure	What it actually counts	Whose number is it	Verifiable?	Safe to say
~80,000	Annual attendance at Burning Man	**Burning Man's.** The host event.	Yes — publicly documented	✅ "Burning Man gathers roughly 80,000 people each year — the environment this work is rooted in."
100 staff / 1,000 volunteers	Burning Man's operating ratio	Burning Man's	Yes	✅ As on `/manifesto`, verbatim
~500	The figure one internal document gives for an actual Oracle deployment	**Playa AI's**	Internal only, unsourced	⚠️ Only with the source and definition attached
80+	Connect pilot users	Playa AI's	portfolio of small shipped pilots	✅ "Connect — a shipped pilot with 80+ users"
~80	Ikigai, Menta Series Miami workshops	Playa AI's	same	✅ "approximately 80 participants"
358	OpenClaw hatchathon community members	Playa AI's	same; flagged as entering the record on one drafting day	⚠️ Define it first: cumulative attendees across N sessions since when
~200	Builders/artists/researchers/funders meeting fortnightly	Playa AI's	same	✅ With the cadence stated
≈520+	Correct aggregate of Playa AI's own participants across pilots and workshops	Playa AI's	derived	✅ With the definition attached

The distinction, stated once, plainly

Burning Man's 80,000 attendees is a true fact about the venue. Playa AI's own participant counts are in the **hundreds**. Both statements can appear in the same paragraph. What must never happen is the first number appearing where the second belongs.

What is safe to say, and what is not

❌ Do not write	✅ Write instead
"derived from 80,000 consented participants in a high-complexity field environment (Burning Man)"	"developed in a high-complexity field environment — Burning Man, an annual gathering of roughly 80,000 people — and tested in Playa AI's own pilots with 80+ participants"
"field-tested at 80,000 participants"	"field-tested in the operating conditions of an 80,000-person gathering, across pilots with 80+ participants each"
"a field study at Burning Man 2024 (80,000+ opt-in participants)"	"field work at Burning Man 2024, an event of 80,000+ attendees; Playa AI's own consented cohorts number in the hundreds"
"80,000+ consented research participants"	"80+ consented Connect pilot users; ~80 Ikigai workshop participants; ~520 across all pilots and workshops to date"
"EARLY PARTICIPANTS 80+ ACROSS PILOTS AND WORKSHOPS" (understates)	"≈520 people across two pilots, a workshop series and monthly build sessions, 2025–26"

Note on the recommended fix for outbound material. Every one of these is a **rewrite**, not a deletion. The 80,000 figure earns its place in all of them — as the description of where the work happens. It just cannot be the answer to "how many participants do you have?"

The single structural defence is finding 2's fix: give the correct sentence a permanent URL, so no future draft has to reconstruct it.

---

5. Links and CTAs

Per-route measurements, captured 2026-07-27

Route	Links	Gmail-compose	`mailto:`	Buttons	Body chars
`/`	32	12	0	5	12,713
`/work`	20	7	0	2	13,396
`/events`	18	2	0	0	6,373
`/team`	11	2	0	0	2,355
`/eelp`	10	2	0	0	4,305

Route	Links	Gmail-compose	`mailto:`	Buttons	Body chars
`/living-trust-framework`	12	3	0	0	12,156
`/manifesto`	8	1	0	0	7,132
Total	**111**	**29**	**0**	**7**	—

Bundle-level count: **45** Gmail-compose URLs (the extra 16 are on routes and states not rendered in a single pass, e.g. `/privacy-toolkit`, `/classic/*` and modal contents).

The seven doors, and where they land

Extracting the `su=` parameter from every Gmail link on the homepage: twelve links, eleven distinct subjects, **one recipient**.

CTA label	Subject line sent
Support Year One	`Donor Inquiry`
Schedule donor conversation	`Donor Inquiry`
Request the brief	`Donor Brief Request`
Request field report updates	`Field Report Notification`
Join the research	`Human Connection Research`
Join the working group	`Working Group Interest`
Join the artist interest list	`Artist Grant Interest` ⚠️
Propose a collaboration	`Research Collaboration`
Start the conversation	`Cultural Stewardship`
Request adopter conversation	`Founding Adopter Inquiry`
Join the initiative	`Join the Playa AI Initiative`
Partner with us	`Partner with Playa AI`
hello@playa-ai.org	`General Inquiry`

No body prefill, no routing, no owner named, no response-time commitment.

Differentiated doors converging on one inbox is not automatically bad design — a single intake with good triage is a legitimate pattern. What makes this a false promise is that **the promise is made in the door copy and abandoned at the threshold.** "*Become a Founding Adopter… Test EELP in real community*

settings. The Founding Adopter pathway is open" describes a pathway; the pathway is an empty compose window. And in one case the subject line commits harder than its own label: "Join the artist interest list" sends `Artist Grant Interest`, asserting a program the same page calls "in development."

Broken and misleading link states

Issue	Detail
`/privacy`	Renders the client-side 404. The **live Connect app links to it** as "Read our privacy policy".
Soft-404	Every path returns HTTP 200, so broken links do not signal as broken to any crawler or link checker.
Self-canonical 404	`/privacy` emits `canonical: https://playa-ai.org/privacy` while displaying "404 · PAGE NOT FOUND" — an explicit instruction to index a not-found page.
7 button-not-link CTAs	Manifesto CTAs and the pledge are `<button>` with no href — not crawlable, not right-clickable, not openable in a new tab, no URL to share. 5 of the 7 are on the homepage; their destinations genuinely exist (`/manifesto` routes correctly and sets the title "Burner-AI Manifesto — Playa AI"), so the href is available and simply not exposed.
Unsignable pledge	Modal headed "FOUNDING SIGNATORIES / PUBLIC COMMITMENT", ten statements, 0 forms, 0 inputs, no signatory list.
Mystic Arts Foundation	Named in the footer of all seven pages, never a link.
No community surface	Across all seven pages there is no link to any social account, newsletter, Discord, Slack or community platform. The complete external inventory is nine event-specific Luma URLs, three GitHub links to one repo, `connect.playai.party`, two prototype hosts, and 29 Gmail links.
No Luma calendar link	Nine per-event Luma links, no subscribable calendar.

6. Copy quality — before and after

The prose on this site is better than most nonprofit copy and the problems are concentrated in specific, replaceable sentences. Each rewrite below preserves the meaning and removes the overreach, the staleness or the register break.

1. The participant stat (understates; finding 10)

> **Before:** EARLY PARTICIPANTS / 80+ / ACROSS PILOTS AND WORKSHOPS
>
> **After:** PARTICIPANTS TO DATE / ≈520 / ACROSS TWO PILOTS, A WORKSHOP SERIES AND MONTHLY BUILD SESSIONS, 2025–26

2. The tools tally (finding 5, finding 15)

> **Before:** Connect App released. Ikigai App piloted in Miami. Agent Labs running monthly. Privacy Toolkit in review.
>
> **After:** Connect App — live pilot, 80+ users. Ikigai App — piloted in Miami with ~80 participants. Agent Labs — last session March 2026. Living Trust Framework — public working draft, licence pending.

3. The events header (finding 12)

> **Before:** PUBLIC EVENT ARCHIVE · SALONS / HACKATHONS / FIELD GATHERINGS · UPDATED CONTINUOUSLY
>
> **After:** PUBLIC EVENT ARCHIVE · SALONS / HACKATHONS / FIELD GATHERINGS · LAST UPDATED 30 MAR 2026 · NEXT SESSION TBA — join the list to hear first

4. The prohibited hero line (finding 14)

> **Before:** Born on the Playa. Gifted to the Planet.
>
> **After:** Built in the desert. Published for everyone.
>
> (Or simply delete the line — the headline above it already carries the page.)

5. The `/events` opening paragraph (finding 14)

> **Before:** Playa AI events are live laboratories for builders, founders, researchers, artists, and operators exploring human-AI co-evolution with intention.
>
> **After:** Playa AI events are working sessions for builders, founders, researchers, artists and operators. Each one produces something checkable: a prototype, a protocol revision, or a published finding.

6. The manifesto programme list (finding 13)

> **Before:** *[section heading]* What We're Building — THE ORACLE / THE FIRST AI CAMP / AI FOR HUMANITY / THE AI CODE OF THE PLAYA / A DIGITAL BURNING MAN
>
> **After:** *[section heading]* What we intend to build — These are directions, not current programs. For what is actually running today, see Work.

7. The Founding Adopter door (finding 7, section 5)

> **Before:** Become a Founding Adopter… Test EELP in real community settings. The Founding Adopter pathway is open. → *[Gmail compose]*
>
> **After:** Become a Founding Adopter. You get the full EELP v1.0 text, a facilitator briefing, and a named contact. We ask for one written field report within 90 days. We reply within five working days. → *[/get-involved#adopters, a real form]*

8. The microgrant ask (finding 9)

> **Before:** Fund the standards work, governance design, Connect pilots, artist microgrants, documentation, and public field reporting.
>
> **After:** Fund the standards work, governance design, Connect pilots, documentation and public field reporting. *(Restore "artist microgrants" when the program has eligibility, a budget and a review process published.)*

9. The manifesto date stamp (finding 13)

> **Before:** MANIFESTO · PLAYA AI · PUBLISHED 2026
>
> **After:** MANIFESTO · PLAYA AI · FIRST PUBLISHED 2025 · LIVING DOCUMENT, LAST REVISED 2026

10. The origin block that does not yet exist (finding 2)

> **New, at `/about#origin`, lifted from `/manifesto`:**
>
> **Where this work is rooted.** Burning Man is a temporary city in the Black Rock Desert. Every year, 100 staff and 1,000 volunteers guide 80,000 people into a living experiment in radical self-reliance, gifting, decommodification, civic responsibility, and creative expression. Playa AI is a separate initiative that works in that environment. Our own pilots and workshops have reached approximately 520 people to date.

That last one is the highest-leverage paragraph in this document. It states the true fact about the host event, states Playa AI's own numbers as Playa AI's own, and gives both a permanent URL to be copied from.

---

7. Tone and voice

The register problem

The site speaks in three registers and switches between them without warning.

1. **Governance register** — `/work`, `/eelp`, `/living-trust-framework`. Hedged, numbered, status-badged, willing to publish open problems. This is the site's best voice and the one that earns institutional trust.
2. **Cultural register** — `/manifesto`, the homepage manifesto block, the footer's "BLACK ROCK DESERT ORIGIN". Incantatory, second-person-plural, unhedged. "*We do not wait for permission. We do not fear the future—we build it.*"
3. **Startup-marketing register** — the stats strips, the seven doors, "The rooms are the product."

The org has already diagnosed this and written the rule down. Prohibited phrases for public-facing website copy names the failure mode exactly: language that "*reads as a spiritual TED talk and causes skeptical institutional audiences (Davos delegates, Stanford researchers) to bounce.*" Its test is the best sentence in the whole vault on this subject: "*If it sounds like a TED talk → rewrite. If it sounds like a governance document with a soul → ship it.*"

The register break is not that register 2 exists. It is that register 2 is unlabelled and unbounded. A manifesto is allowed to be a manifesto. The problem is (a) two prohibited phrases from register 2 leaked into the homepage hero and `/events`, and (b) `/manifesto` mixes register-2 prose with an unlabelled list of five unbuilt programs, so a reader cannot tell where the poetry stops and the roadmap starts.

Who is alienated

Reader	What loses them
Program officer at an institutional funder	The hero line "Born on the Playa. Gifted to the Planet." before any evidence; "human-AI co-evolution" on `/events`; a pledge with no signatories; five unbuilt programs in the present progressive.
Academic researcher	No methodology anywhere; no definitions behind any headcount; the EELP gated behind an email; 358 published with no denominator.
Diligence officer / grants compliance	`/privacy` 404s next to an app that records audio; no LICENSE on the "open-source" flagship; "artist microgrants" advertised with no program; the fiscal sponsor unlinked.
Engineer or builder considering adoption	A framework that cannot legally be forked; two prototype apps on Replit and Manus scratch hosts; no link to Connect from the page describing Connect.
Burner / community member	Almost nothing — this is the one audience the current voice serves well. But there is no community surface at all: no Discord, no newsletter, no calendar, no upcoming anything.
Screen reader or reduced-motion user	46 unstoppable scroll animations; 54.8% of homepage text at opacity 0 until scrolled; a modal that traps the virtual cursor without focus.

The voice it should have

One voice, in two clearly labelled rooms.

- **The default across `/`, `/work`, `/events`, `/team`, `/eelp`, `/living-trust-framework`:** the governance register the org already writes well. Concrete nouns, dated claims, status badges, numbers with definitions, open problems stated before achievements. This voice is genuinely uncommon in this sector and it is the org's actual competitive advantage in a funder's inbox.
- **`/manifesto`, clearly framed as what it is:** a founding cultural text. Give it a dateline, a one-line frame at the top ("This is our founding cultural text, first published 2025. For what is running today, see Work"), and let it be as incantatory as it wants inside those walls. Its own in-100-words note already does this well — "*Note: This is Playa AI's founding cultural text. Its current operating expression is being developed through EELP, the Living Trust Framework, community experiments, and published learning.*" That sentence is the model. It should be at the top, not buried mid-page.

The single rule that resolves everything: **every register-2 sentence must be inside `/manifesto`, and every register-2 sentence must be adjacent to a link to a register-1 page.**

---

8. Structure and information architecture

Page weight

Route	Body chars	Links	In nav?	Assessment
`/work`	13,396	20	Yes	Strongest content, worst navigation within it (finding 21)
`/`	12,713	32	Yes	Doing eight jobs; competing portfolio with `/work`
`/living-trust-framework`	12,156	12	**No**	Substantial and honest; unreachable from nav
`/manifesto`	7,132	8	**No**	Zero inbound links from anywhere (finding 2)
`/events`	6,373	18	Yes	Real evidence; four months stale, archive-only
`/eelp`	4,305	10	**No**	Most conservative page on the site; off-nav
`/team`	2,355	11	Yes	Thinnest page; every name listed twice
`/about`	—	—	—	**Does not exist**; blocked on Founding Story approval
`/get-involved`	—	—	—	**Does not exist**; scoped in March 2026, never built
`/privacy`	—	—	—	**Does not exist**; linked from the live Connect app
`/privacy-toolkit`	live	—	No	Undocumented, stale, indexable
`/classic/*`	live	—	No	Undocumented legacy surface, 4 routes

Four of the seven public pages are off-nav. Three pages that the architecture *requires* do not exist. Five undocumented routes are live and indexable.

The seven-doors problem

The homepage presents a seven-audience matrix — builders, artists, researchers, funders, cultural stewards, institutions, general — with a differentiated promise for each. All seven converge on one inbox differentiated only by an email subject line (section 5). This is the clearest structural symptom of `/get-involved` never being built: **the CTA architecture was designed for a landing page that does not exist**, and the Gmail links are what filled the hole.

Either make the doors real (each becomes a section of `/get-involved` with criteria, a form and a stated next step) or collapse them into one honest door: "Email hello@playa-ai.org and tell us which of these you are." Seven promises and one inbox is the worst of both.

Recommended navigation

```
Playa AI

  Work          → /work            (canonical portfolio, 8 workstreams, index at the top)
  Standards     → /standards       (hub: EELP, Living Trust Framework, Data Trust design)
  Events        → /events          (Upcoming block first, then archive, then Luma calendar)
  About         → /about           (origin + the 80,000 block + team + council + fiscal sponsor)
  Manifesto     → /manifesto       (framed as founding cultural text)
                                    ─────
  [Get Involved] → /get-involved   (primary CTA, persistent, present on mobile)
```

Notes on the change: - `/manifesto` **enters the nav.** It is the page that states the numbers correctly and it currently has zero inbound links. - `/standards` **is a new hub** so `/eelp` and `/living-trust-framework` stop being orphans. - `/about` **absorbs** `/team` and carries the origin block, the fiscal sponsor with EIN and link, and the Council Charter. - **One primary CTA**, which is what the 2025-03-13 decision already decided and the shipped site never implemented. - `MMXXVI` **is deleted**; `Contact` folds into `/get-involved#contact` and lives inside `<nav>` so it survives below 768px. - Add a mobile menu. Publish `sitemap.xml`. Add a footer "Next" strip on every page.

---

9. Accessibility, SEO and machine readability

Grade: D+. The content and copy discipline are considerably better than the delivery. Nearly every defect in this section is a configuration or one-file fix rather than a rewrite.

The mission-level argument

The organisation's own stated principle, on its own homepage, is **Gifting**:

> "*Open-source the standards. Share the learning publicly. The work becomes infrastructure when the lessons travel freely.*"

A standard cannot travel through machines that cannot read it. Every route serves an identical 1,717-byte shell whose body is `<div id="root"></div>`. Googlebot, GPTBot and `facebookexternalhit` were each tested directly and each received the same empty shell. There is no prerendering tier. Every non-JS consumer — social link unfurlers, AI search retrievers, archive tools, link checkers, and the automated diligence tooling funders increasingly run — sees a blank page with homepage boilerplate.

This lands hardest on the numbers. **The website has already done the right thing:** it states Burning Man's 80,000 attendees as Burning Man's, and its own participant counts separately and modestly. But that correct framing exists only in executed JavaScript, with no SSR, no structured data and no sitemap — while the conflated framing ("derived from 80,000 consented participants") sits in outbound PDFs and a policy-database submission as plain, perfectly indexable text.

The org fixed the page and left the machine-readable record uncorrected. Any AI or diligence tool asking "how many participants does Playa AI have?" retrieves the conflation and cannot reach the correction. Fixing rendering is therefore not a cosmetic SEO chore — it is the mechanism by which the accurate framing becomes reachable at all.

Findings summary

Area	Finding	Severity
Performance	No compression at all; 668 KB → 132 KB with brotli; `cache-control: private` on hashed immutable assets	**critical**
Machine readability	No SSR/prerender for any consumer including Googlebot and GPTBot	**critical**
Technical SEO	One hardcoded `og:url` and `og:title` for the whole site; no subpage has a social identity	**critical**
Technical SEO	Soft-404s returning HTTP 200; the 404 view self-canonicalises; no valid `sitemap.xml`; no `Sitemap:` in robots.txt; no JSON-LD anywhere	high
Trust	`/privacy` 404s while the live Connect app links to it; stale `/privacy-toolkit` archive is the top indexable candidate for "Playa AI privacy"	**critical**
Mission alignment	No LICENSE on the framework repo; owner is a User account, not an Organization	**critical**
Accessibility	`prefers-reduced-motion` ignored at both CSS and JS layers; 46 `whileInView` animations; 54.8% of homepage text at opacity < 0.1 until scrolled	high
Accessibility	Pledge modal never moves focus into itself; 37 background focusables remain tabbable; no `inert`, no focus trap	high
Accessibility	All images carry `alt=""`	medium
Conversion	45 Gmail-compose links, 0 `mailto:`, 0 forms, 0 inputs	**critical**
Colour contrast	Exactly **one** genuine WCAG AA failure site-wide after resolving `oklab()` colours correctly	low

A correction to our own method

An earlier pass reported ten contrast failures and several unlabelled inputs and stray `a.csm-action` links. Nine of the ten contrast failures were artifacts of mis-parsing `oklab()` colour values; after compositing alpha through a canvas, only one real failure remains. The stray links, "Post" submit buttons and four

unlabelled inputs were **injected by a Diigo browser extension in the test profile**, not site code, and have been excluded. Similarly, "7 CTAs are `<button>` without href" is correct as a whole-site count but resolves to 5 on the homepage specifically. Reporting these accurately matters more than a bigger number.

---

10. Drift — what the org corrected after the corpus closed

Nothing in this table is an error. These are places where the site is *newer* than the brain, and in most cases the change is the org moving toward more accuracy. They are separated so they are never read as findings, and the credit is due plainly.

What changed	Vault state at corpus close	Site now	Read
The 80,000 conflation is gone from the website	Detected as a critical unresolved conflict; the number appears 315 times across 65 Mission Control items	The number appears **only** on `/manifesto`, correctly attributed to Burning Man with the operating ratio stated; Playa AI's own figure is stated separately as "80+ across pilots and workshops"	**The single best result in this audit, and it deserves explicit credit.** The org did not delete a fact — it *un-conflated* it. The true statement about the host event was kept; the borrowed participant count was replaced with countable pilots. Corroborating: the figure appears **zero** times across all 21 July 2026 grant chats even though Burning Man is discussed in 19 of them, and in the same session the team ran an explicit scrub on a sibling claim (drop the "millions" audience claim: "*Naked 'audience: millions' is the line that makes a reviewer distrust everything above it*"). The remaining exposure is off-site (finding 2), not on the page.
"Foundation" → "initiative"	Rebrand next-action, open, logged as an eligibility risk on the OpenAI application	"*a human-centered AI initiative, part of Mystic Arts Foundation's California 501(c)(3)*"	Executed after corpus close. Correct move. Residual: MAF's formal governance of the initiative is still unconfirmed, and "public charity" as a sub-classification is not established anywhere in the vault.
A public EELP draft page exists	Open question: publish a public v1.0 draft before institutional review?, open; squad rule to reference v1.0 externally only as "targeted for Q2 2026"	`/eelp` live, labelled DRAFT, "*not yet finalized, certified, or legally approved*"	Answered affirmatively after the corpus closed, and labelled *below* what the vault would support. The most conservative page on the site. Q2 ended 2026-06-30, so July is a modest slip, not a broken commitment.
Living Trust Framework repo is public	Publication unverified; repo missing LICENSE, CHANGELOG, CONTRIBUTING	Public repo created 2026-07-15, with a candid README that self-polices its own licence status	The repo appearing is drift, and the README's honesty is a genuine credit. The missing LICENSE is not drift — see finding 4.

What changed	Vault state at corpus close	Site now	Read
Ikigai App 1 CTA now reads "CURRENTLY UNAVAILABLE"	Featured to the OpenAI Foundation as one of two items with real participant counts	Dead CTA, honestly labelled	Downward, good-faith disclosure — the site labels it rather than hiding it, which is the behaviour §08 promises. Sequencing is unflattering (eleven days after it was named as evidence), but the label is right.
9 events; 8 workstreams	Six convenings recorded; eight work-page items confirmed 2026-07-16	9 / 8	Workstreams: exact match, confirmed. Events: the count is plausible drift; only the roster is a finding (11).
`/playa-ai-living-trust-framework` alias handled correctly	Not recorded	Serves the same page and canonicalises to `/living-trust-framework`	Someone did the SEO thinking properly here. Credit where due.
Mission and positioning language	Various earlier framings	Current framing	The org is entitled to evolve how it describes itself. No findings manufactured here.

11. What the site is missing that the record shows it has earned

The brain holds better material than the site publishes. In rough order of credibility gained per unit of effort:

1. **The crisis protocol.** EELP §6.1 hard-wired into the Oracle flow: on language indicating acute distress, the system exits its normal flow at any phase, delivers a fixed script referencing the 988 Lifeline, notifies a human facilitator — "*this is not optional and cannot be handled by the AI alone*" — and generates no artifact for that session. Two sentences convert "human override" from a slogan into a testable safety behaviour. **The highest credibility-per-sentence asset in the vault, and it is on no public surface.**
2. **The event log with outcomes.** The hackathon series records a 174-minute session on 2026-03-04 where participants generated OAuth tokens, provisioned VPS instances and stood up Telegram bots, plus a CRM reconciliation of 46 registered / 19 attended. `/events` currently republishes Luma marketing copy with zero results. Prerequisite: the corpus carries four irreconcilable attendance figures for Hackathon 1 — reconcile to one number with a definition, or publish none.
3. **What Agent Labs actually offers.** Subsidised VPS hosting at roughly $27/month capped at the first 25 signups, a spend-visible LLM proxy layer, and a deliberate security posture of running agents on an isolated remote host rather than the member's own machine (community agent onboarding). That is a concrete offer; "work toward deploying a live agent" is not.
4. **The fortnightly community call.** Twelve distinct sessions from 2025-04-09 to 2026-03-25 — the longest-running thread in the corpus — with a recurring roster of named researchers and funders. The site has no recurring-community surface at all. This is the most durable proof of an actual community the org owns. (Count twelve meetings, not sixteen files; Fireflies double-captured.)

5. **The Data Trust's actual architecture.** A five-seat Trust Council with named composition, 4-of-5 quorum, four-layer independently revocable consent, a Public Use Registry, and a stated deletion window (data trust and Trust Council; participants hold the data governance). Publishing the *design* while the body is unconstituted is honest and specific. Rule 13 already supplies the language discipline for doing it without implying enforceability.
6. **What is actually inside the EELP** — v1.0 plus fourteen numbered addenda, §7.2 no-affective-outcome claims, §8.1 facilitator training as a compliance requirement, §10.0 a tiered certification ladder, §5.1b enacted consent via the Carry Card, and TRACE, a six-step anonymisation pipeline (EELP artifact). A numbered clause list is the difference between a standard and a poster — and it directly tests the site's own claim "*Designed to be adopted without a legal team.*"
7. **The Oracle's engineering and governance record.** Oracle, Burning Man 2026 holds a six-phase interaction flow marked governance-cleared 2026-03-07, biometrics explicitly excluded from the MVP, and matchmaking deliberately rejected as untested. The site names "Oracle" once, in passing, and otherwise only in the manifesto's aspirational register. Publish it as a design standard, not as a thing standing in the desert — the form factor was still unconfirmed at corpus close and the July grant chats never mention it.
8. **"On Our Own Terms" and "Ask First"** — a designed twelve-month capacity-building programme (30 artists, 10 cultural organisations, 12 workshops, an openly licensed consent toolkit) that would give the empty "Fund" pillar something to point at (portfolio of small shipped pilots). Publish as the Year-One design, not as funded.
9. **The Council Charter.** Already written, already good, and the exact thing a diligence officer looks for. `/team` publishes ten names and no governance document.

Vault gaps this audit exposed — things the brain should hold and does not: no Ikigai artifact or project page despite it being one of two headline pilots; no decision record stating the Burning-Man-attendance versus Playa-AI-participants distinction; no record of whether the March funder sends went out; no record of the live site as shipped (two competing March copy packages, no decision between them, no launch record); no event log or definition of "Playa AI public event"; no participant-overlap or aggregate figure; no sourcing methodology behind any headcount; no microgrant program of any kind; no artifact page for a "Privacy Toolkit"; and no "Year One" program page, though the site uses the name three times.

---

12. Prioritised fix list

Today — under an hour each, do all of them

#	Fix	Finding	Effort
1	**Replace all 45 Gmail-compose URLs with** `mailto:hello@playa-ai.org?subject=…`	3	30 min
2	**Enable brotli/gzip; set** `cache-control: public, max-age=31536000, immutable` **on** `/assets/*`	17	30 min

#	Fix	Finding	Effort
3	**Add the LICENSE file to the framework repo**	4	30 min once the licence is chosen
4	**Delete "Privacy Toolkit in review" from the homepage**	5	1 line
5	**Wrap the app root in `<MotionConfig reducedMotion="user">` + CSS backstop**	16	30 min
6	**Remove "Born on the Playa. Gifted to the Planet." (×2) and "human-AI co-evolution" (×1)**	14	20 min
7	**Convert the 7 manifesto `<button>`s to `<a href="/manifesto">`; delete the modal**	2	45 min
8	**Change the artist Gmail subject from `Artist Grant Interest` to `Artist Interest List`**	9	2 min
9	**Delete the `MMXXVI` span**	18	2 min
10	**Remove "artist microgrants" from the donor CTA until the program exists**	9	2 min
11	**Fix the manifesto year stamp (2026 → first published 2025)**	13	2 min
12	**`noindex` on `/privacy-toolkit` and `/classic/*`**	1, 6	15 min

Total: about four hours of work, and it removes three of the six critical findings.

This week

#	Fix	Finding	Effort
13	**Ship `/privacy`** — what is collected, where it goes, retention, deletion contact; point the Connect app at it	1	half a day
14	**Rewrite the Connect disclosure screen** — microphone, recording, upload, transcription, retention, deletion, before the button	1	1 screen of copy
15	**Prerender the seven routes** with per-route `<title>`, description, `og:url`, `og:image`	6	1–2 days
16	**Real HTTP 404s; remove the self-canonical; publish `sitemap.xml` + `Sitemap:` in robots.txt**	6	3 hours
17	**Fix the participant aggregate** — "≈520 across two pilots, a workshop series and monthly build sessions, 2025–26"	10	30 min + defining 358
18	**Add the origin block** at `/about#origin` with the Burning Man paragraph, permanently linkable	2	half a day

#	Fix	Finding	Effort
19	**Add focus management to the pledge modal**	16	1 hour
20	**"Upcoming" block on `/events`; "Last updated" stamps site-wide; Luma calendar link**	12	2 hours
21	**Reconcile the portfolio to one canonical list of eight**	15	1 hour
22	**Label `/manifesto`'s "What We're Building" as intent; add a CTA at its foot**	13, 19	30 min
23	**Publish `/llms.txt`** stating Playa AI's own participant figures and the Burning Man distinction	6	1 hour
24	**Add a mobile menu; move `Contact` inside `<nav>`**	18	2 hours

Before the next funder conversation

#	Fix	Finding	Effort
25	**Rewrite the three outbound vectors carrying the conflated phrasing** — the OECD.AI / Partnership on AI submission first (highest exposure, hardest to unwind), then the Mozilla email v1.2 whose subject line *is* the conflation, then confirm which March funder emails went out and issue corrections. **Rewrite for accuracy, do not retract — the Burning Man context is true and belongs in all three.**	2	1–2 hours per document
26	**File the decision record** in the vault stating the Burning-Man-attendance vs Playa-AI-participants distinction, so it cannot recur through a draft	2	20 min
27	**Correct the grant record** with the OpenAI Foundation: the framework is a public working draft, licence pending — unless (3) has already made it true	4	one email
28	**Build `/get-involved`** with seven anchored sections and a real form	7	1–2 days
29	**Make the pledge signable** and render the signatory list	8	1 day
30	**Publish the nine-row event log** with London restored, plus a definition of "Playa AI public event"	11	half a day
31	**Build a real `/about`** absorbing `/team`: origin, fiscal sponsorship with MAF linked and EIN stated, legal structure, people with affiliations, Council Charter	20	1 day + Founding Story approval
32	**Close the MAF regranting question in writing**	9	an email and a wait
33	**Publish the crisis protocol** — the highest credibility-per-sentence asset the org owns	§11	1 hour
34	**Rebuild `/work` §01 as an actual index; move "What remains unresolved" to the top**	21	2 hours

#	Fix	Finding	Effort
35	**Consolidate product domains under** `playa-ai.org`	22	a few hours
36	**Add the prohibited-phrase list to the repo and grep it in CI**	14	1 hour

13. What is genuinely good

This is not performative balance. These are real strengths and several of them are uncommon enough to be competitive advantages.

1. **The org un-conflated the numbers itself, before anyone made it.** The website states Burning Man's 80,000 attendees accurately, attributed to Burning Man, with the operating ratio spelled out — and states Playa AI's own participation separately and modestly. Most organisations in this position keep the bigger number and hope. This one kept the true fact and gave up the borrowed credential. That is the behaviour a funder is actually testing for, and it happened without external pressure.
2. **Status discipline on** `/work` **is the real thing.** Eight workstreams, each with a status tier. Unprompted hedges like "*Technical practices listed here are implementation directions, not fully deployed capabilities.*" An entire section headed "*OPEN QUESTIONS / What remains unresolved*" with six numbered items including security testing and anti-sycophancy. Very few early-stage orgs publish their own unsolved problems.
3. **The** `/eelp` **and** `/living-trust-framework` **pages label themselves *below* what the vault would support.** `/eelp` says "*not yet finalized, certified, or legally approved.*" The Living Trust Framework says "*The open-source license is under review.*" Both are more conservative than what the org told the OpenAI Foundation eleven days earlier. The website is the honest surface.
4. **The Living Trust Framework README is unusually self-policing.** Verbatim: "*License status: Under review. No open-source or content license is currently granted unless and until an approved* `LICENSE` *file is added.*" It also flags that the document "*includes historical AI-generated draft material, proposed practices, aspirational governance, and placeholders that are not current Playa AI commitments*", and preserves a clearly-marked "Historical License Proposal — Not in Force". Nobody is overclaiming here.
5. **The Ikigai App 1 CTA reads "CURRENTLY UNAVAILABLE" rather than hiding the dead link.** Downward disclosure on a headline pilot, eleven days after it was named as grant evidence. That is expensive honesty.
6. **Heading hierarchy is genuinely correct across every route.** Exactly one H1 per page, no skipped levels — `/work` runs H1 → H2 → H3 across 27 headings with no violations. This is the single most common SPA accessibility failure and this site does not have it.
7. **Client-side meta management is well built.** Per-route canonical and `og:url` are correct once JS runs, and the alias route `/playa-ai-living-trust-framework` correctly canonicalises to `/living-trust-framework` rather than creating a duplicate. Someone thought about this properly. The only problem is that it runs too late for non-JS consumers.

8. **The pledge modal's ARIA scaffolding is right:** `role="dialog"`, `aria-modal="true"`, a resolving `aria-labelledby`, an Escape handler, a close button with `aria-label`, and body scroll locking. Only focus management is missing — the bones are correct.

9. **Colour contrast very nearly passes.** After resolving `oklab()` values correctly, exactly one genuine WCAG AA failure exists site-wide.

10. `/events` **is falsifiable.** Nine dated events with per-event external links. A reader can check every one. That is rarer than it sounds.

11. **Landmarks,** `lang="en"`**, zero positive tabindex values, 8** `:focus-visible` **rules, permissive** `robots.txt`**, HSTS with a two-year max-age and** `includeSubDomains`**, and** `rel="noopener noreferrer"` **on every external link.** The fundamentals are in place.

12. `/manifesto`**'s own in-100-words note** — "*This is Playa AI's founding cultural text. Its current operating expression is being developed through EELP, the Living Trust Framework, community experiments, and published learning.*" — is the single best sentence of self-framing on the site. It should be at the top of the page, not buried mid-scroll.

The honest summary: the writing on this site is ahead of the delivery, and the delivery problems are mostly config files. An organisation with this much status discipline and this little server-side rendering is one week of unglamorous work away from a site that reads as credible to both a program officer and a crawler.

Related

- Claim inventory: `.vault-meta/site-claims.md`
- Honest-numbers discipline in grant writing
- Eligibility risk from "Playa AI Foundation" branding vs fiscal sponsorship
- Participant count: 80,000 vs 500 — read as the first detection of the category conflation, not as a dispute about whether a number is true
- Portfolio of small shipped pilots
- Prohibited phrases for public-facing website copy
- Minimum viable launch scoped to 3 pages
- Single primary CTA (Get Involved), manifesto moved to page bottom
- Fix mobile UX on Playa AI website (hidden buttons, broken hover)
- OpenAI People-First AI Fund application
- Candid annex (local only, not committed): `private/website-audit-2026-07-annex.md`